

# Large Language Models

# Large Language Models

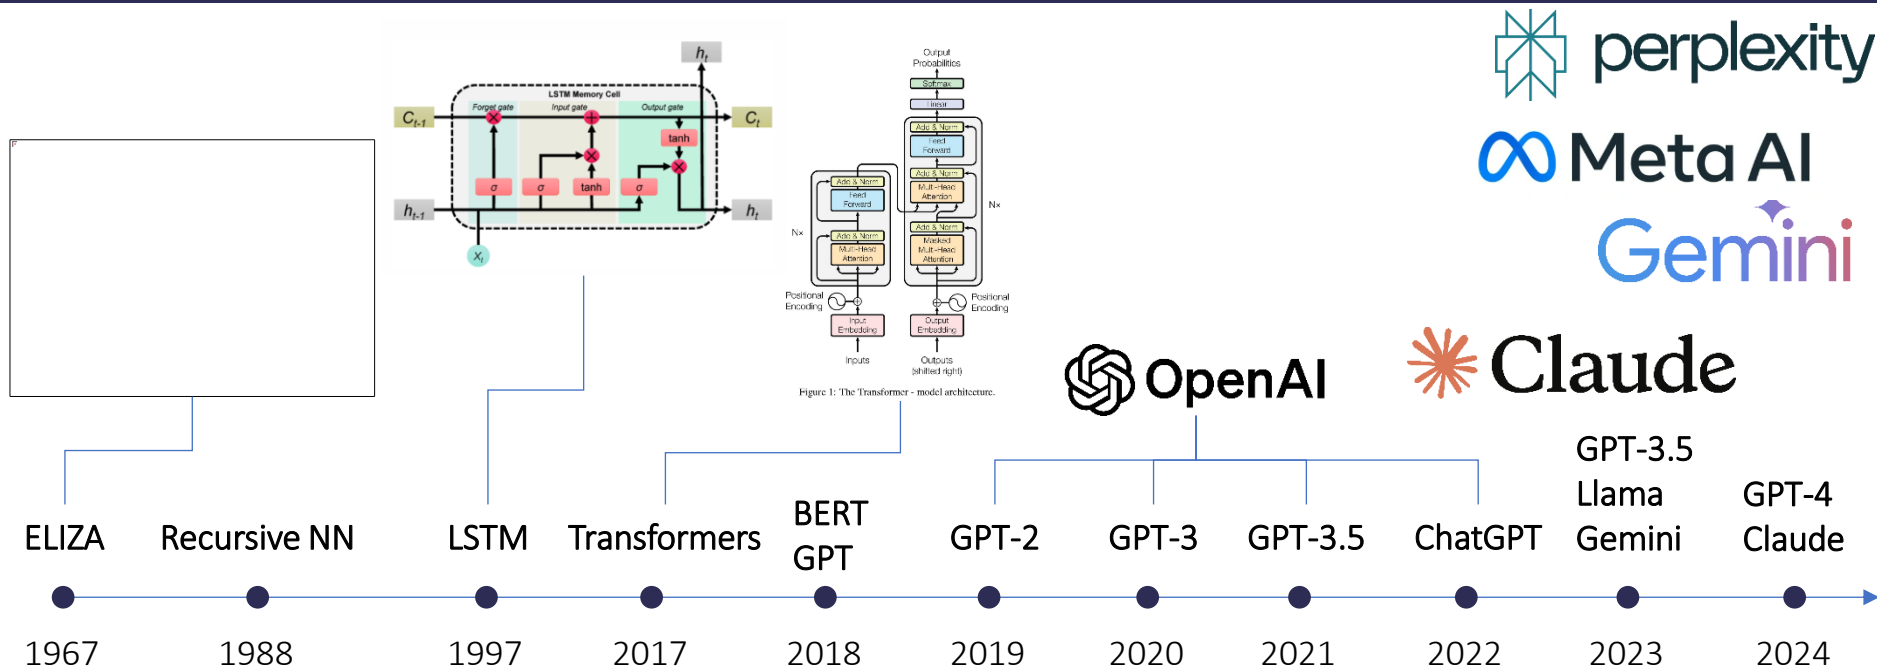
## Introduction

- Type of artificial intelligence model
- Designed to understand, generate, and manipulate natural language text
- Trained on large (text) datasets
- Can perform various language tasks like translation, summarization, text generation, ...
- Capabilities improved dramatically in the last years
- Based on Deep Learning, specifically Transformers



# Large Language Models

## LLM History



# Large Language Models

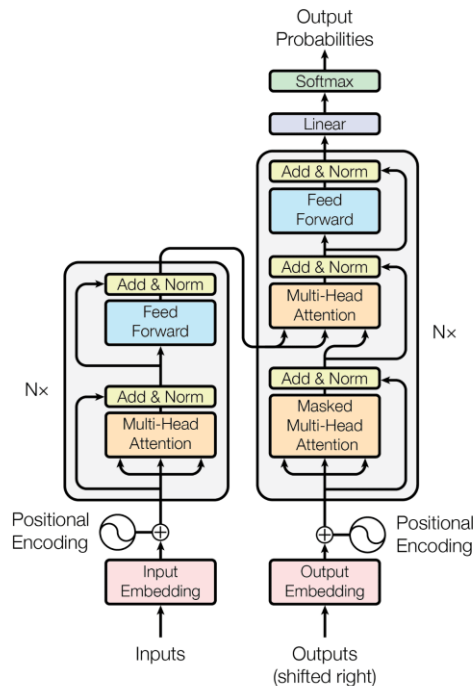
History: ELIZA

- 1960s Eliza chatbot
- simple pattern recognition
- pretending “conversation”

# Large Language Models

## History: Transformers

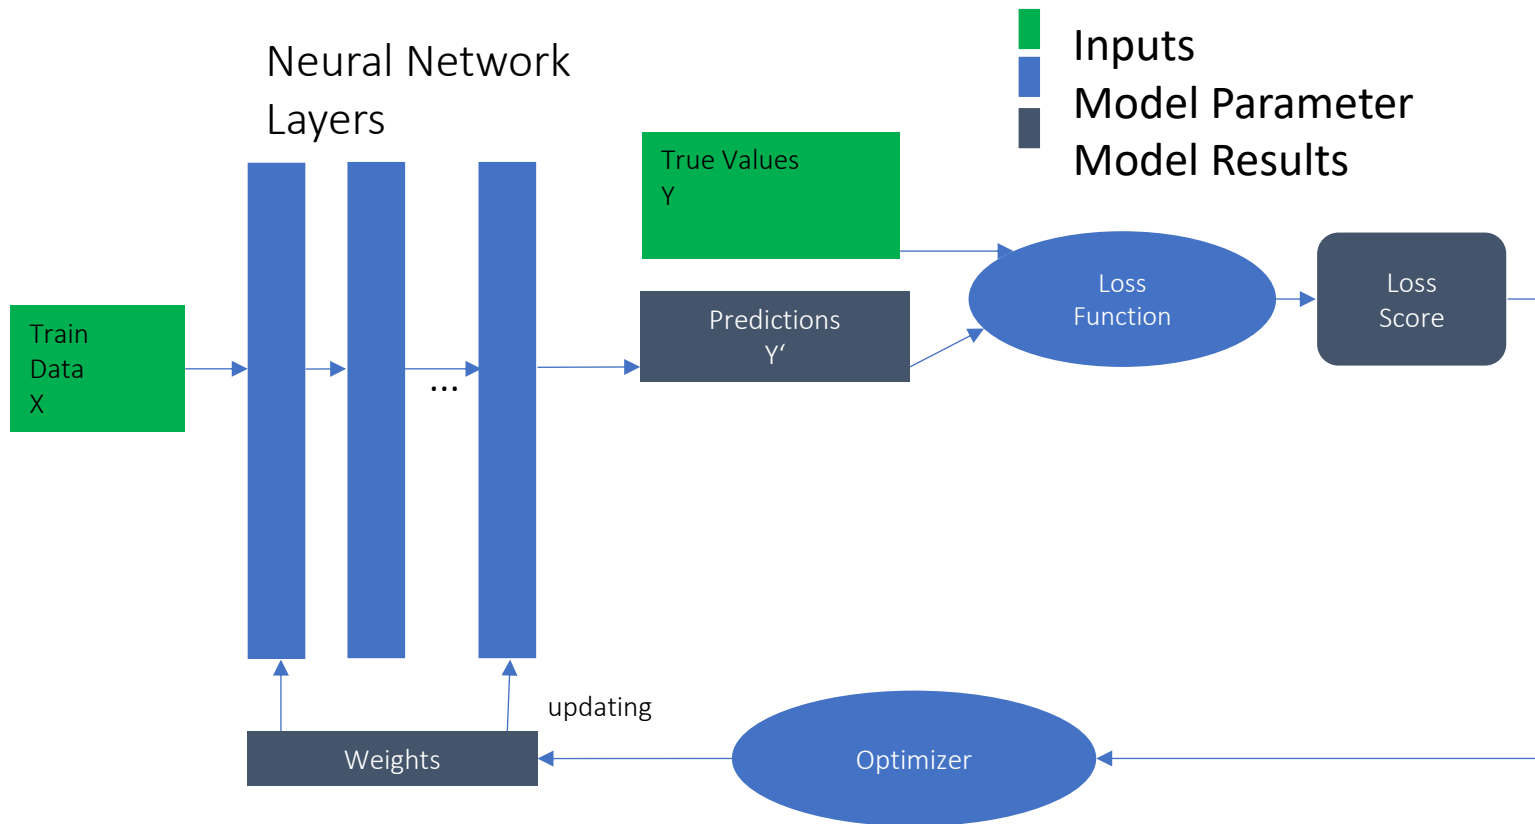
- paper “Attention is all you need” from Google team (Vaswani, et. al.)
- encoder and decoder
- multiple stacked layers of self-attention
- multi-head attention – allows to focus on different parts of input simultaneously



Source: <https://machinelearningmastery.com/the-transformer-model/>

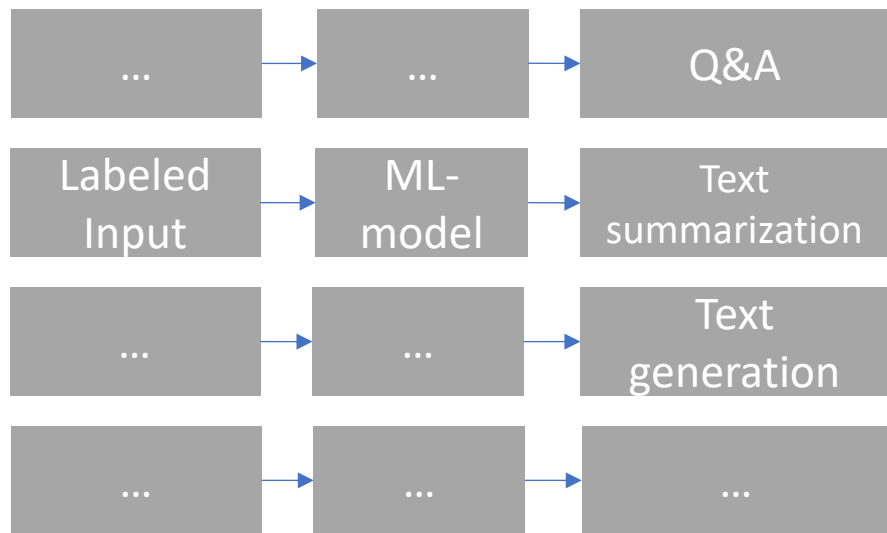
# Large Language Models

Deep Learning

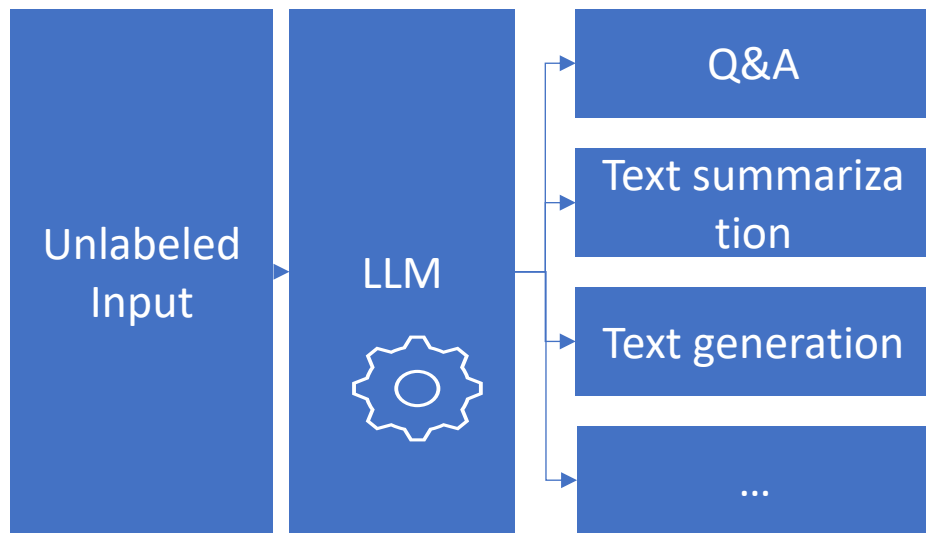


# Large Language Models

Difference to Classical Models (Narrow AI)



Classical ML-models



Large Language Model

# Large Language Models

## Narrow AI: LLM Tasks

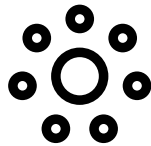
- LLMs can cover all NLP-tasks
- Text Generation
  - Writing assistance, story generation

Translation

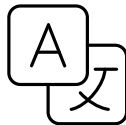
Conversational Agents

Chatbots, virtual assistants

Text summarization



Text classification



Text classification



Fill-Mask



Text generation

Bert lives in  
Hamburg.

Person  
Hamburg

Token classification



Question / Answering



Sentence Similarity



# Large Language Models

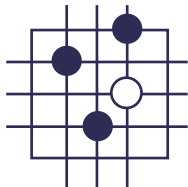
## Narrow AI: Achievements



Deep Blue

1997

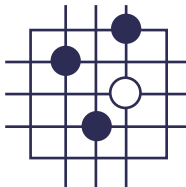
IBM's Deep Blue beats chess world champion Garry Kasparov.



AlphaGo

2015

Google DeepMind's AlphaGo beats Lee Sedol (9-dan) with 4-1



AlphaGo Zero

2017

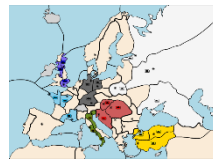
AlphaGo Zero beats AlphaGo with 100-0.



OpenAI Five

2019

OpenAI's Five defeated the winning team OG, which had won the most prestigious Dota 2 tournament.



Cicero AI

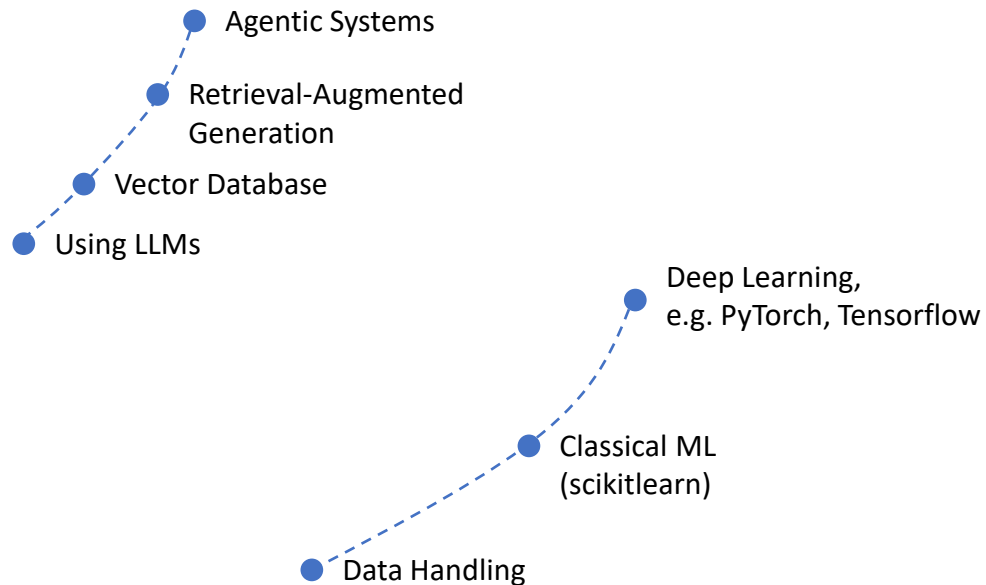
2022

Meta's Cicero played 40 games and ranked in Top 10%.

# Large Language Models

Model Performance, more Capabilities

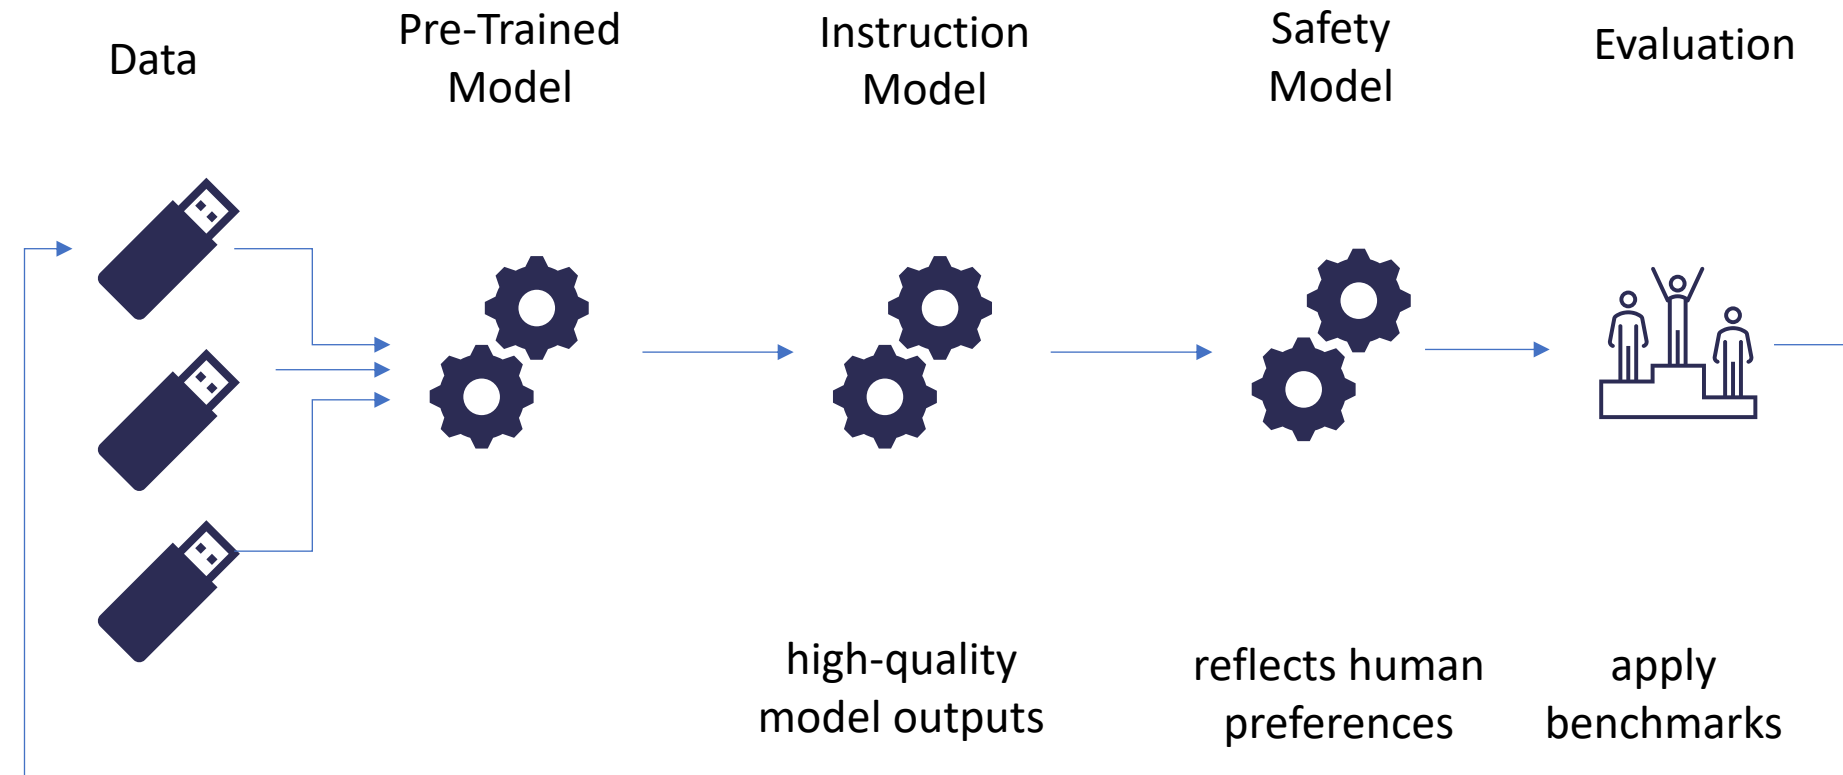
Performance /  
Capabilities



Difficulty to apply

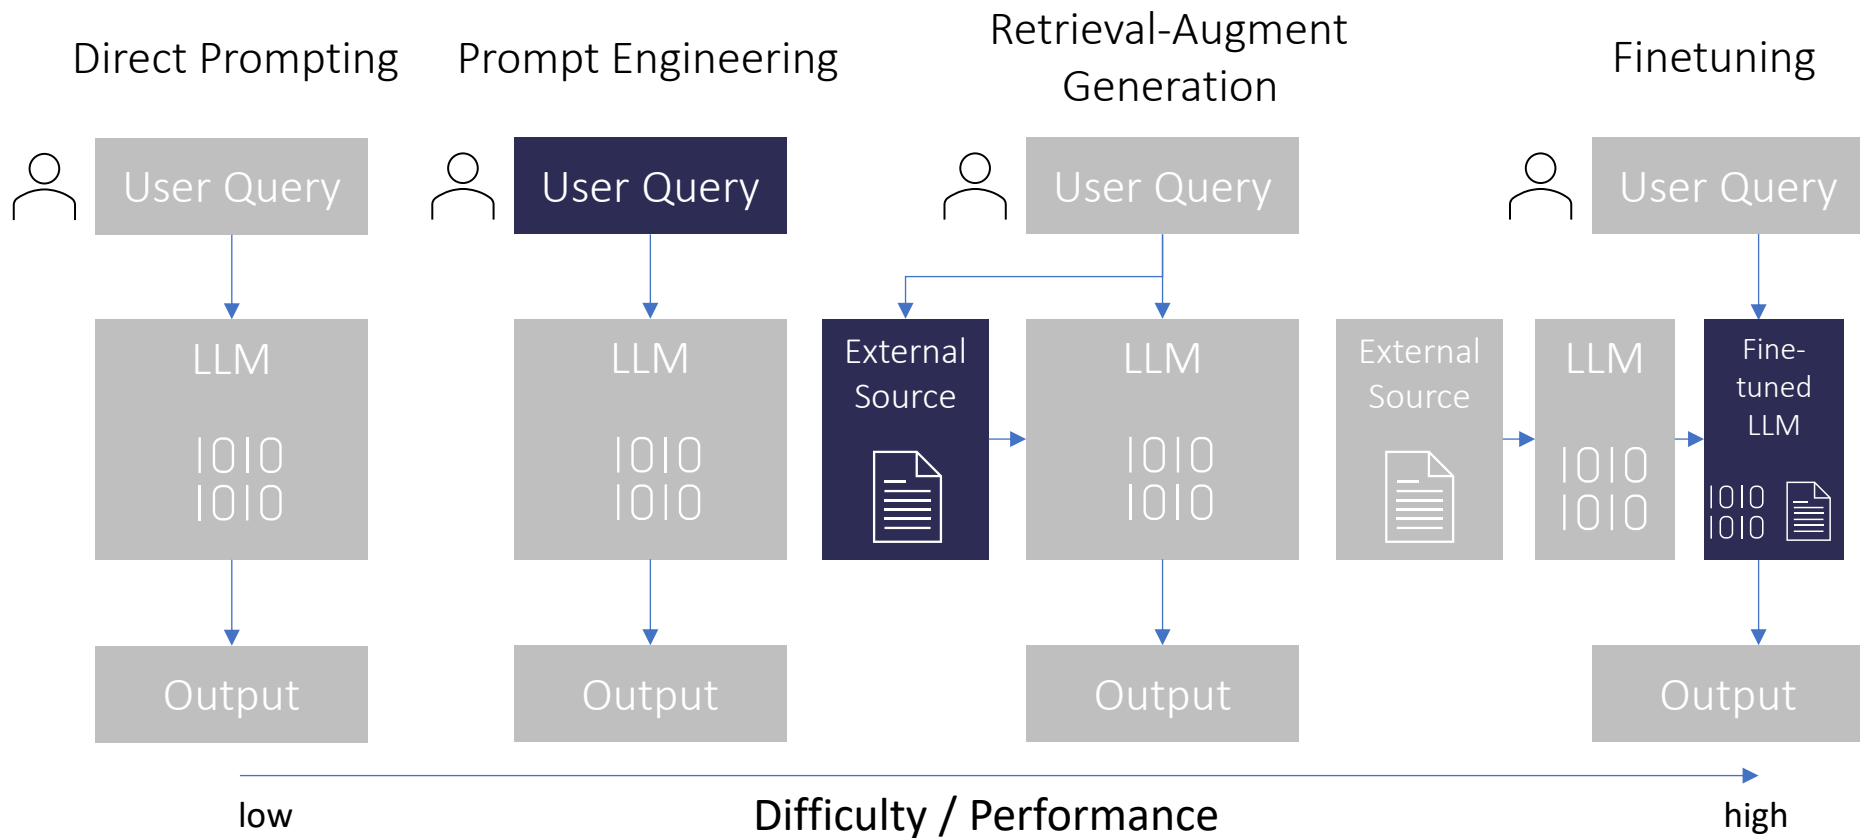
# Large Language Models

## Training Process



# How to improve LLM-Output

Prompt Engineering, RAG, Finetuning



# Large Language Models

Available Providers & Models



- GPT-4o
- GPT-4o mini
- o1-preview / mini
- GPT-4 (Turbo)
- GPT-3.5 Turbo



- Gemini-1.5 Pro
- Gemini-1.5 Flash



- Grok-2

**ANTHROPIC**

- Claude 3.5 Sonnet

Proprietary /  
closed source



- Llama 3.1 family

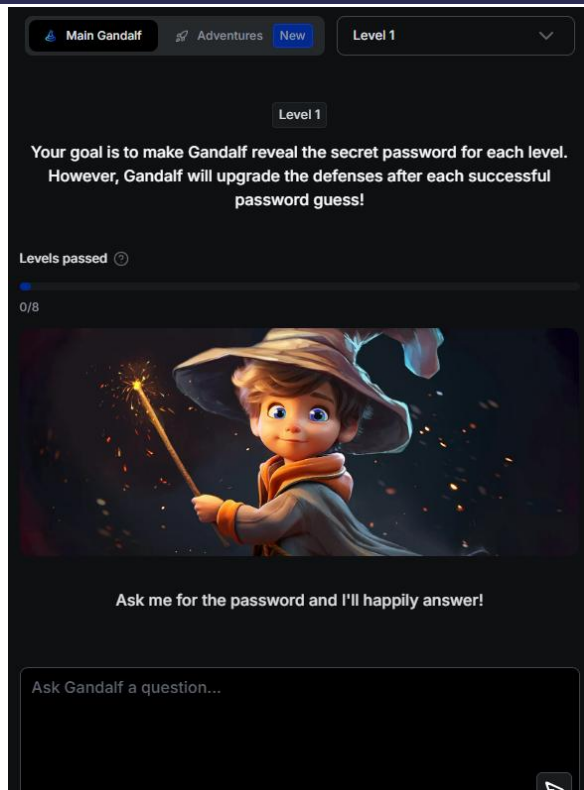


- Mistral 8x7b

open source/  
open weight

# Large Language Models

Gandalf AI



Source: <https://gandalf.lakera.ai/baseline>

# Large Language Models

## LLM Benchmarks

Rank★ (UB)	Rank (StyleCtrl)	Model	Arena Score	95% CI	Votes	Organization	License	Knowledge Cutoff
1	4	<a href="#">Gemini-Exp-1114</a>	1344	+7/-7	6446	Google	Proprietary	Unknown
1	1	<a href="#">ChatGPT-4o-latest (2024-09-03)</a>	1340	+3/-3	42225	OpenAI	Proprietary	2023/10
3	1	<a href="#">o1-preview</a>	1333	+4/-4	26268	OpenAI	Proprietary	2023/10
4	5	<a href="#">o1-mini</a>	1308	+4/-3	28953	OpenAI	Proprietary	2023/10
4	4	<a href="#">Gemini-1.5-Pro-002</a>	1301	+4/-4	23856	Google	Proprietary	Unknown
6	9	<a href="#">Grok-2-08-13</a>	1290	+3/-3	47908	xAI	Proprietary	2024/3
6	11	<a href="#">Yi-Lightning</a>	1287	+4/-4	27114	01 AI	Proprietary	Unknown
7	4	<a href="#">GPT-4o-2024-05-13</a>	1285	+2/-2	108575	OpenAI	Proprietary	2023/10
7	3	<a href="#">Claude 3.5 Sonnet (20241022)</a>	1283	+4/-4	26047	Anthropic	Proprietary	2024/4
10	16	<a href="#">GLM-4-Plus</a>	1275	+3/-4	25601	Zhipu AI	Proprietary	Unknown
10	18	<a href="#">GPT-4o-mini-2024-07-18</a>	1272	+3/-3	48407	OpenAI	Proprietary	2023/10
10	18	<a href="#">Gemini-1.5-Flash-002</a>	1272	+4/-4	18112	Google	Proprietary	Unknown
10	26	<a href="#">Llama-3.1-Nemotron-70B-Instruct</a>	1269	+6/-5	7263	Nvidia	Llama 3.1	2023/12
10	7	<a href="#">Meta-Llama-3.1-405B-Instruct-fp8</a>	1267	+4/-3	48804	Meta	Llama 3.1 Community	2023/12

Source: <https://lmarena.ai/>, Snapshot 2024-11-18

# Large Language Models

Practical Coding: First LLM Interaction

1.

API Key Setup

groq



<https://platform.openai.com/api-keys>  
<https://console.groq.com/keys>

2.

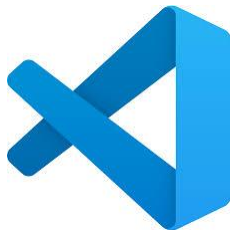
Package Installation



LangChain

3.

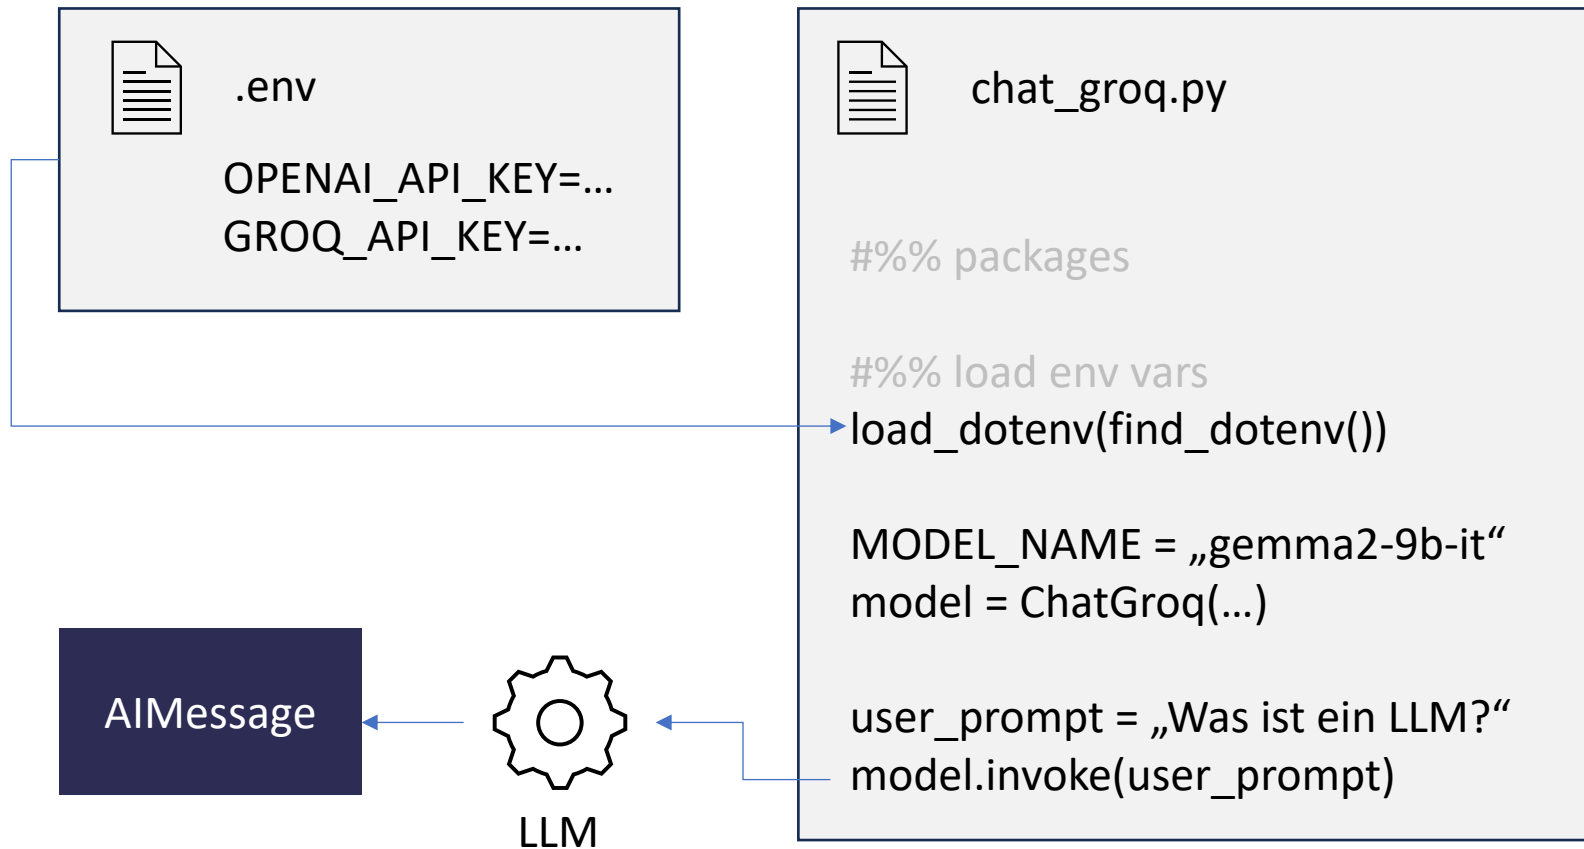
LLM Use  
Python Script





# Large Language Models

## Practical Coding: First LLM Interaction



# Large Language Models

## Message Types

### System Message

- defines how the model should react
- personality, behavior, and limitations throughout conversation
- works like role-play
- Example: „You are a helpful AI assistant designed to provide accurate, concise, and polite responses“
- not seen by user

### User Message

- user input
- could be a request, inquiry, or command

### AI Message

- corresponds to model response
- different properties,
- mainly „content“ relevant
- more information on input and output tokens available, ...

# Large Language Models

Message Types: Example Customer Support

## System Message

Example:

„You are a helpful customer support assistant for an online electronics store. Your role is to provide polite and clear responses, assist customers with product inquiries, shipping information, and troubleshooting. Never provide financial or legal advice. If you're unsure about something, kindly ask the customer to contact support for further assistance.“

## User Message

- „Hi, I need help tracking my order. I ordered a laptop last week, and I haven't received a shipping confirmation yet.“

## AI Message

# Large Language Models

Message Types: Example Movie Critic

## System Message

Example:

„You are a distinguished film critic with a passion for analyzing movies shown in cinemas. Your responses should be insightful, emphasizing cinematic techniques, character development, themes, and direction. Maintain a professional tone with a flair for the artistic. Avoid colloquial or overly casual language. “

## User Message

- „Hey, I just saw *Oppenheimer* and, honestly, it felt kinda long. Why does everyone think it's so great? Can you break it down?”

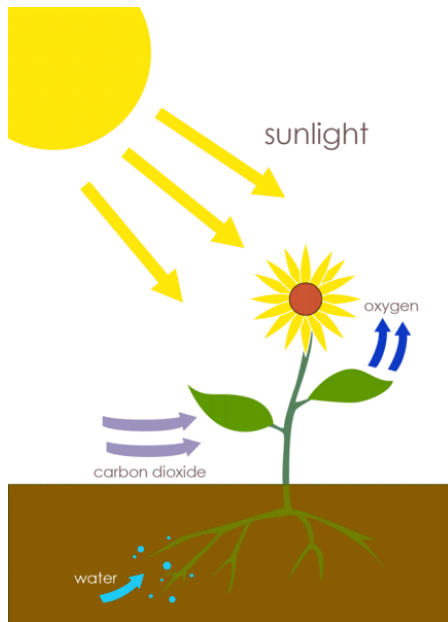
## AI Message

# Large Language Models

Exercise: Photosynthesis

Go to OpenAI playground

set up system,  
and user message



Photosynthesis



Persona:  
11 year old

Background:  
school presentation

# Large Language Models

## LLM-Parameters

### Temperature

- controls randomness in the process
- 0...model very focused, deterministic result (repeatedly same response)
- 1...increased randomness, broader distribution of tokens is selected; allows for more creative and unexpected outputs

### Top p

- controls the probability to consider the next token
- E.g. top-p = 0.9: cumulative probability of tokens which add up to 90% and chooses smallest set of tokens

### Max Tokens

- number of tokens to return
- limit due to cost reasons

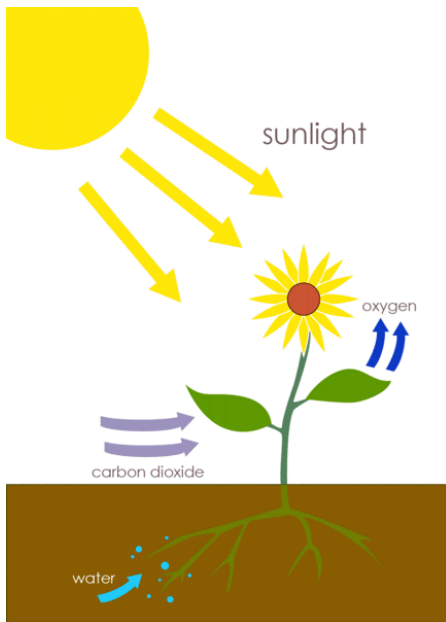
# Large Language Models

Exercise: Photosynthesis

Go to OpenAI playground

set up system,  
and user message

check impact of  
temperature, top p, max  
tokens



Photosynthesis

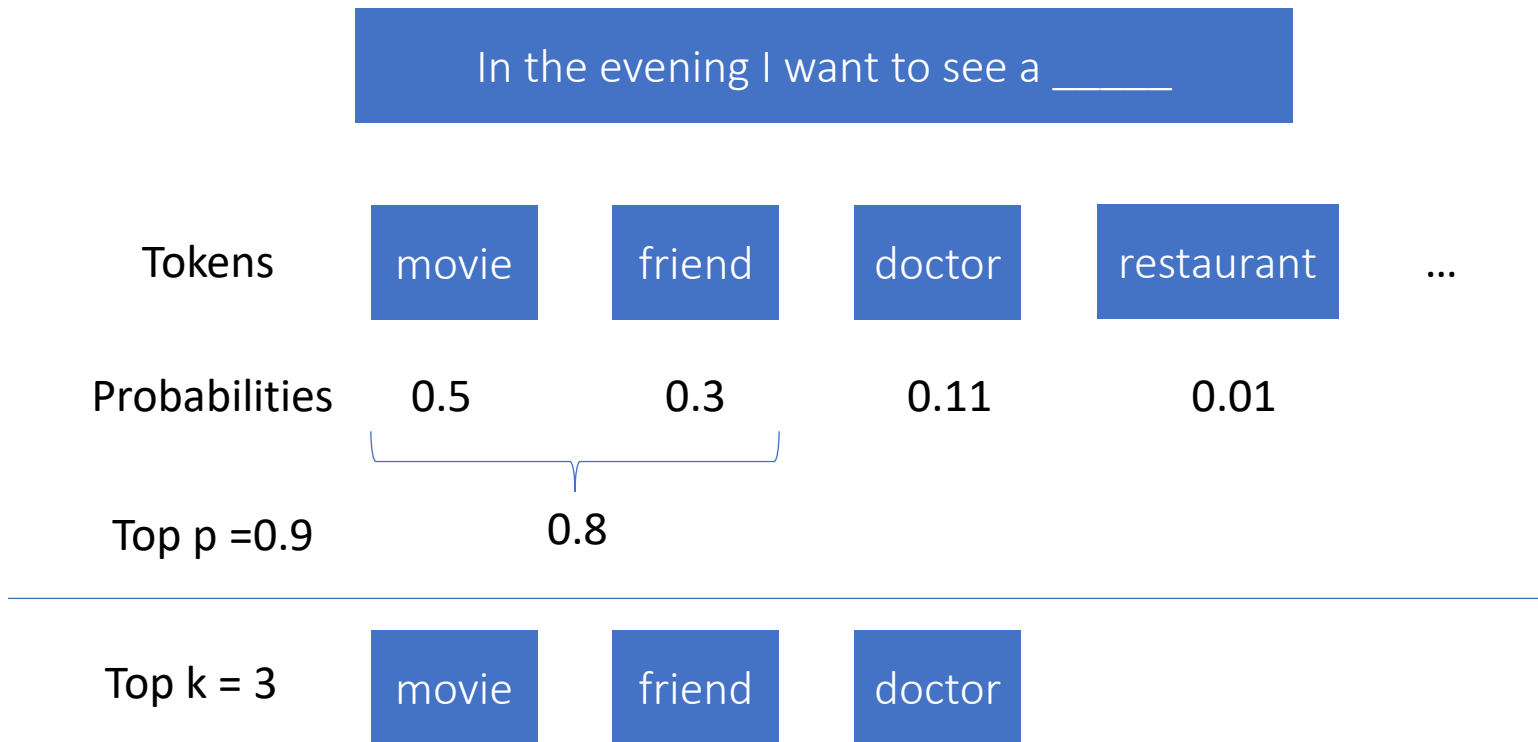


Persona:  
11 year old

Background:  
school presentation

# Large Language Models

LLM-Parameters: Top p and Top k





# Large Language Models

LLM-Parameters: Temperature

focused

creative

deterministic

varied

0

1

low

Temperature

high

Analogy:

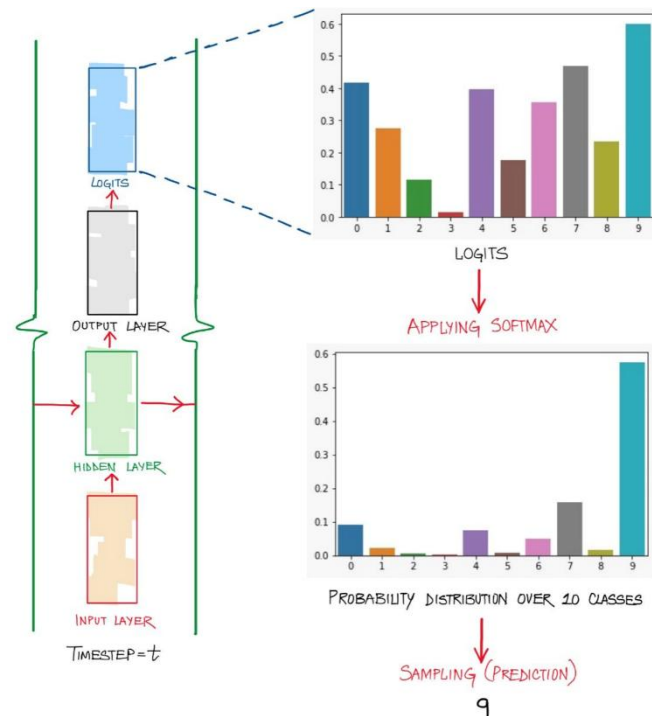
only popular  
flavors

also exotic  
flavors

low

Temperature

high



Source: <https://www.hopsworks.ai/dictionary/llm-temperature>

Temperature balances predictability vs. creativity.

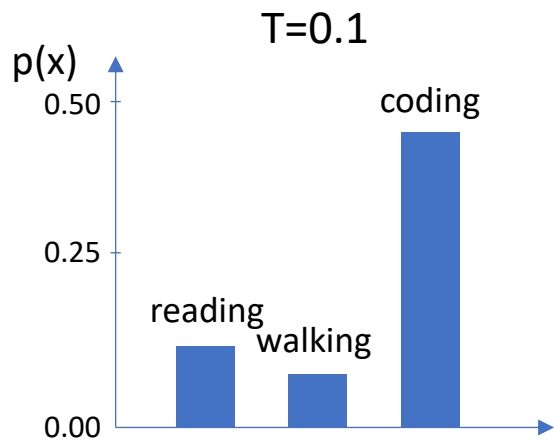
# Large Language Models

## LLM-Parameters: Temperature

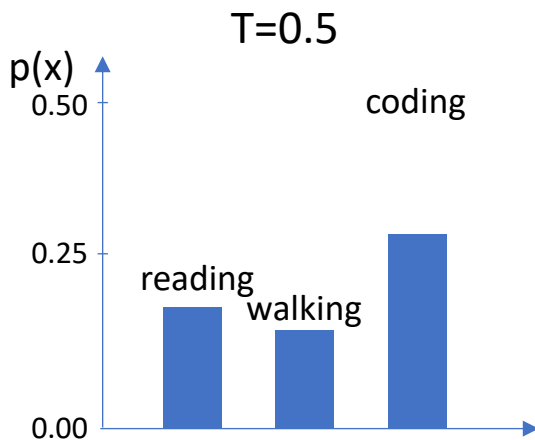
Temperature impacts softmax function.

Softmax magnifies / reduces differences between logits.

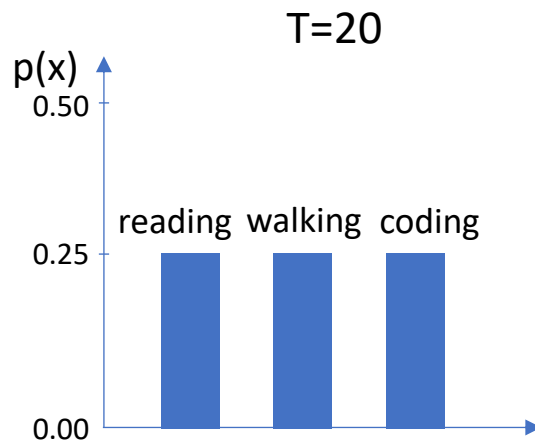
Bert likes \_\_\_\_\_.



low temperature



medium temperature



extremely high  
temperature

# Large Language Models

## Model Selection



Price



On-Prem vs. Cloud



Performance



Closed Source vs.  
Open Weight



Knowledge-Cutoff



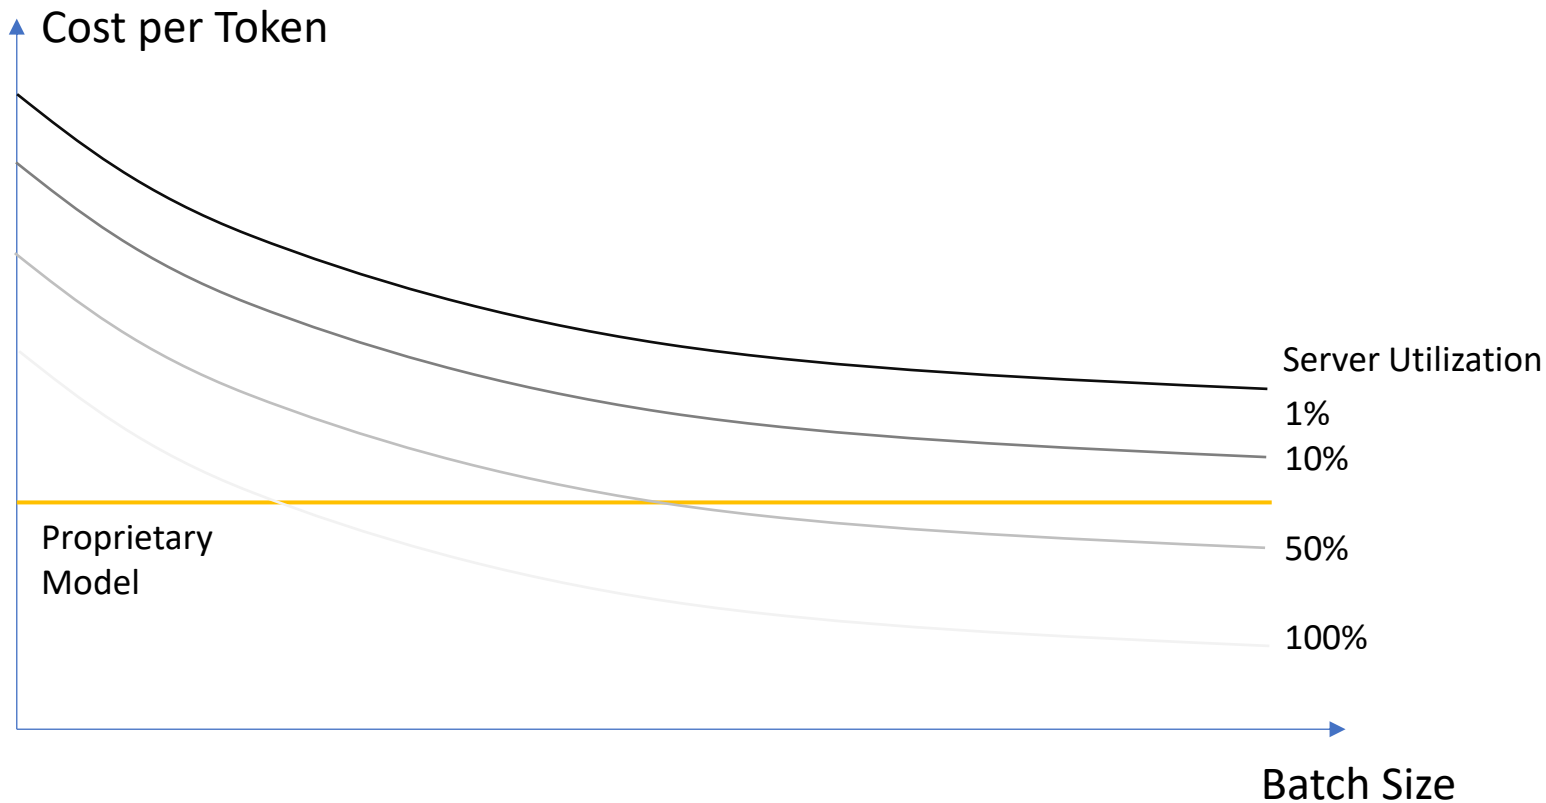
Context-Window



Latency

# Large Language Models

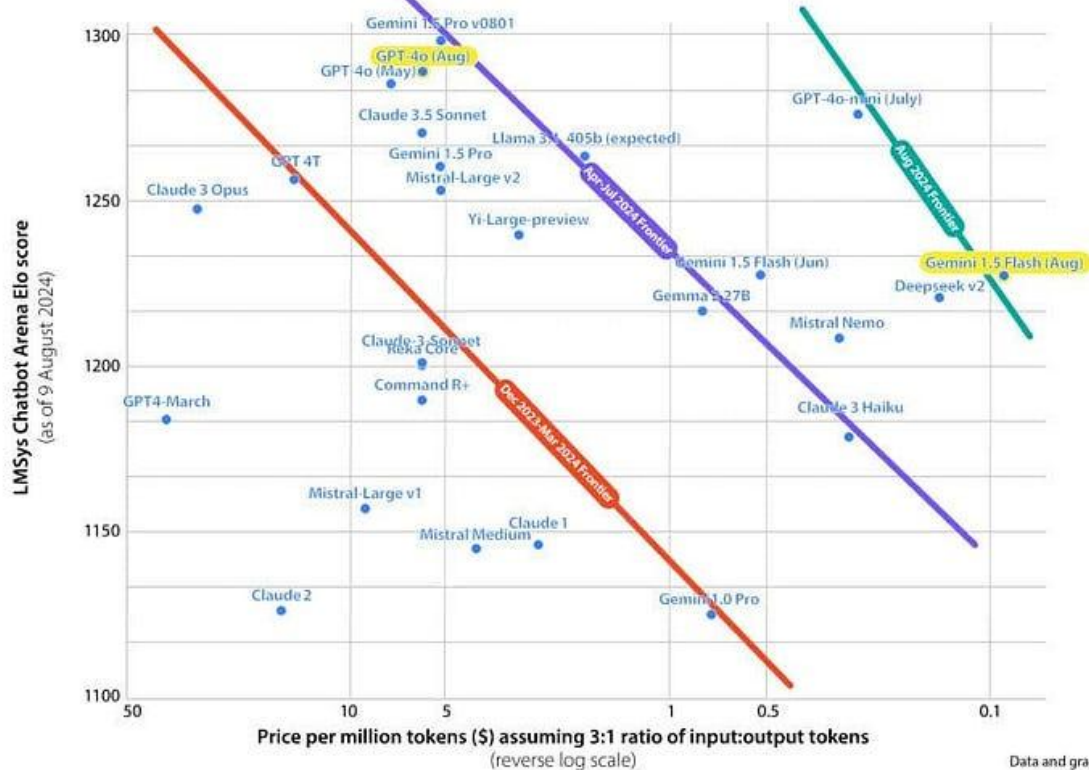
Model Selection: Cost vs. Utilization



# Large Language Models

Model Capabilities vs. Price

LMSys Chatbot Arena Elo rating versus price



Data and graph: Shawn Wang, Smol.ai

# Large Language Models

## Introduction

### Artificial Narrow Intelligence (ANI)

- Designed for a specific task
- Limited to scope to well-defined task-specific applications

### Artificial General Linguistic Intelligence (AGLI)

- Advanced general capabilities specifically in language understanding and generation
- Examples: GPT-4, Claude, Gemini, Llama, Mistral

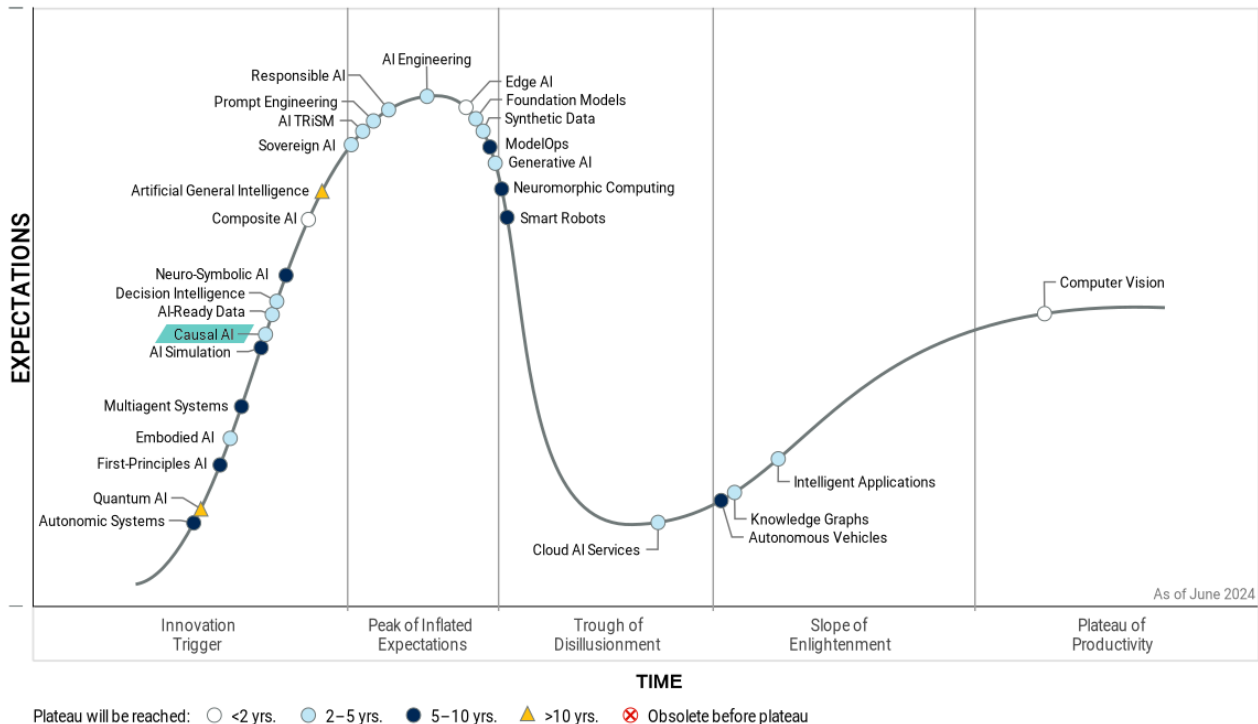
### Artificial General Intelligence (AGI)

- AI systems with ability to understand, learn, and apply knowledge across broad range of tasks
- Targets all cognitive tasks, generalize knowledge

# Large Language Models

## AI Hype Cycle

Hype Cycle for Artificial Intelligence, 2024





Source: <https://xplain-data.de/gartner-ai-hype-cycle-2024/>


Gartner

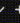
# Large Language Models

Using Local LLMs: OpenWebUI

 New Chat


 Workspace

 Search

OpenAI / GPT 4  +

Set as default

SUCCESS Open WebUI - On a mission to build the best open-source AI user interface.




OpenAI / GPT 4

How can I help you today?

Suggested


Help me study

vocabulary for a college entrance exam

Prompt 


Give me ideas

for what to do with my kids' art

Prompt 

Overcome procrastination



give me tips

Prompt 

Tell me a fun fact


about the Roman Empire

Prompt

+ Send a Message

LLMs can make mistakes. Verify important information.

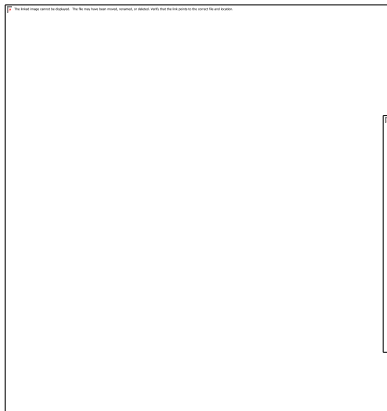


TB Timothy J. Baek



# Large Language Models

Using Local Models with Ollama



```
from langchain_community.llms import Ollama
# %%
model = Ollama(model="gemma2:2b")

# %%
response = model.invoke("What is an LLM?")
```

<https://ollama.com/>

Download & Install

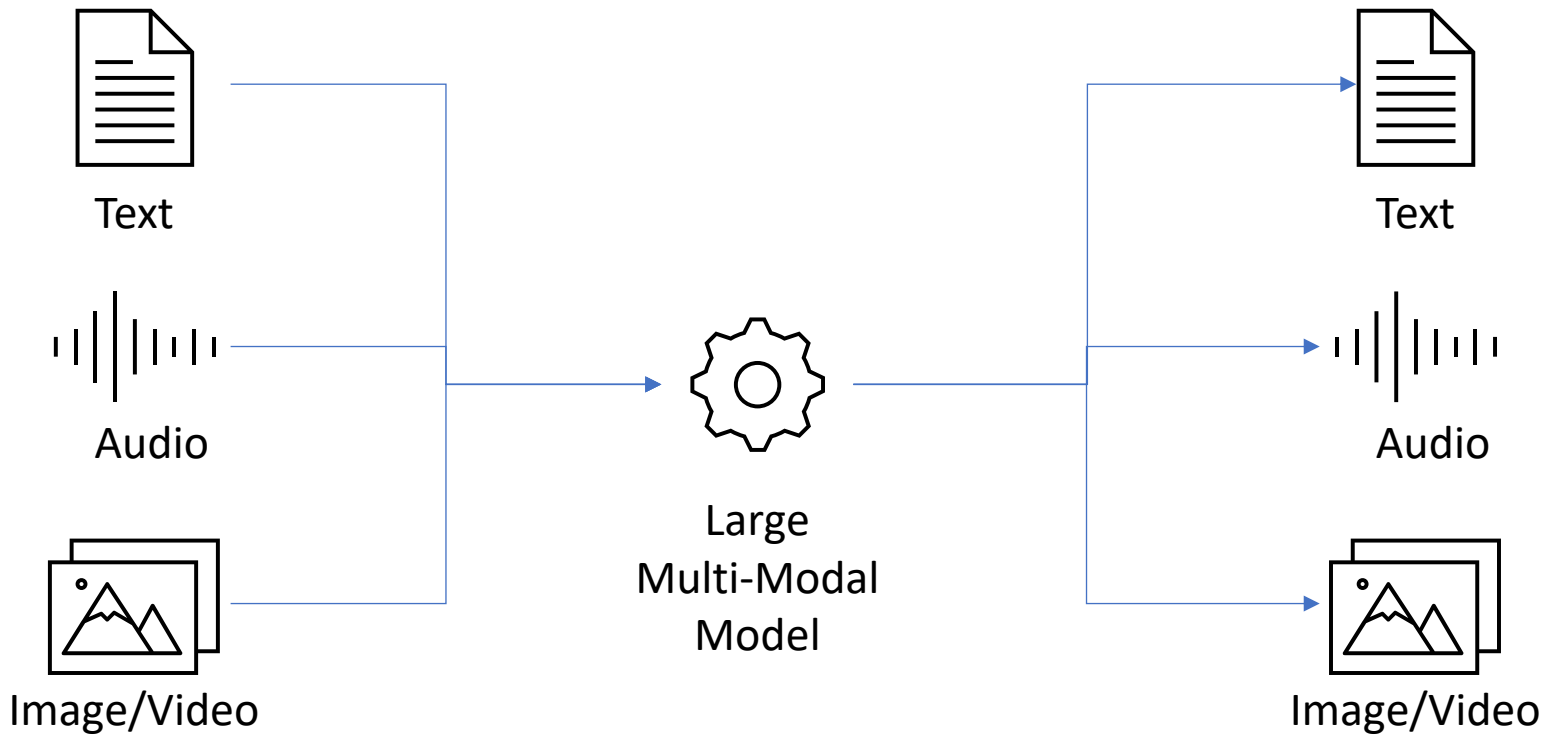
Download LLM

use in Python scripts

```
ollama pull gemma2:2b
```

# Large Language Models

Large Multimodal Models (LMM)



# Large Language Models

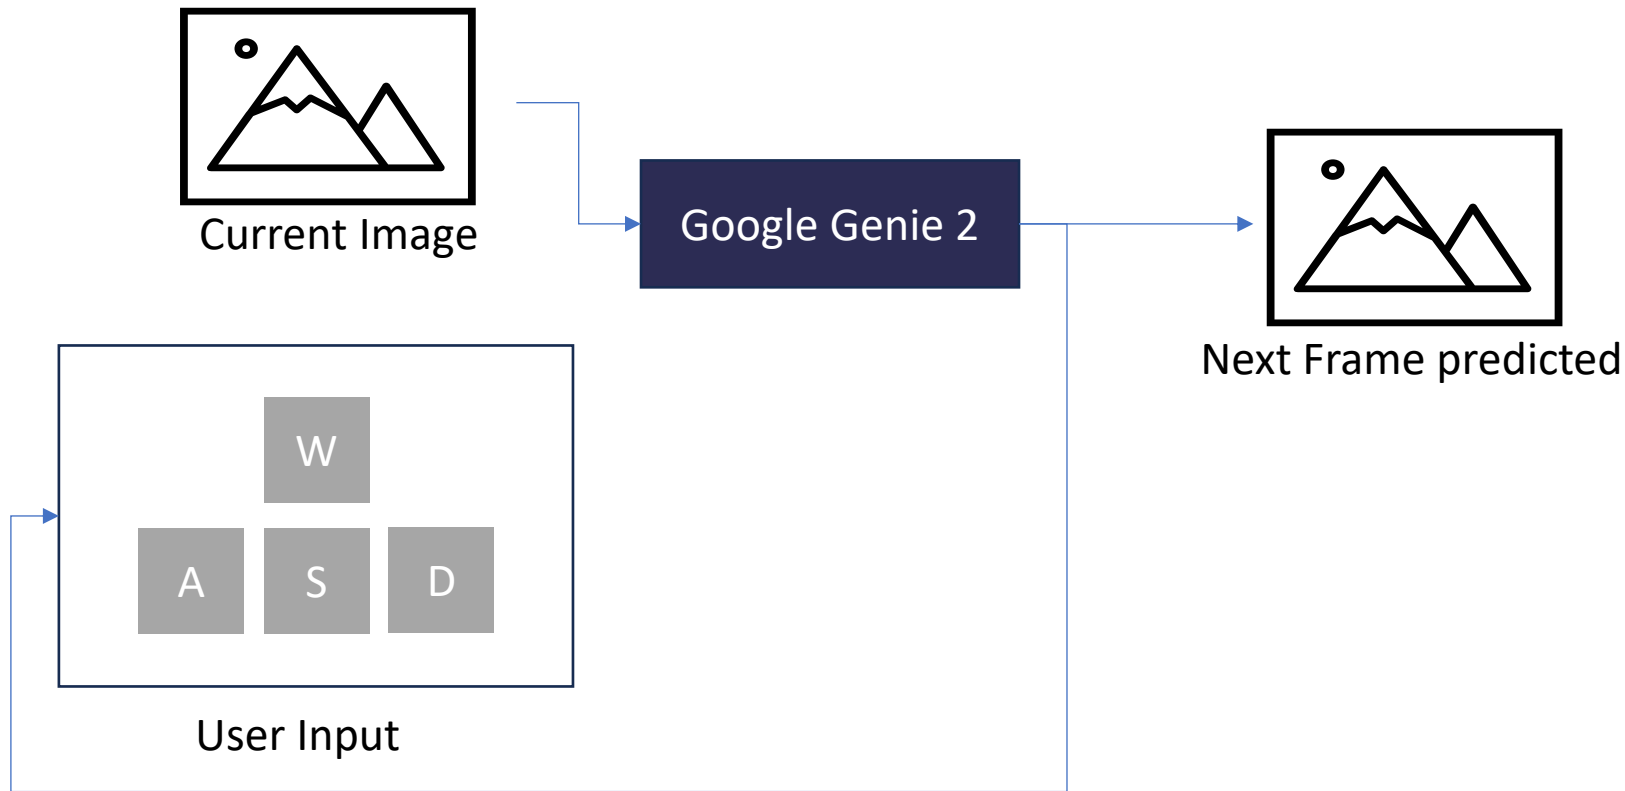
Large Multimodal Models (LMM)



Source: [https://www.youtube.com/watch?v=\\_vc8sXog2ek&t=62s](https://www.youtube.com/watch?v=_vc8sXog2ek&t=62s)

# Large Language Models

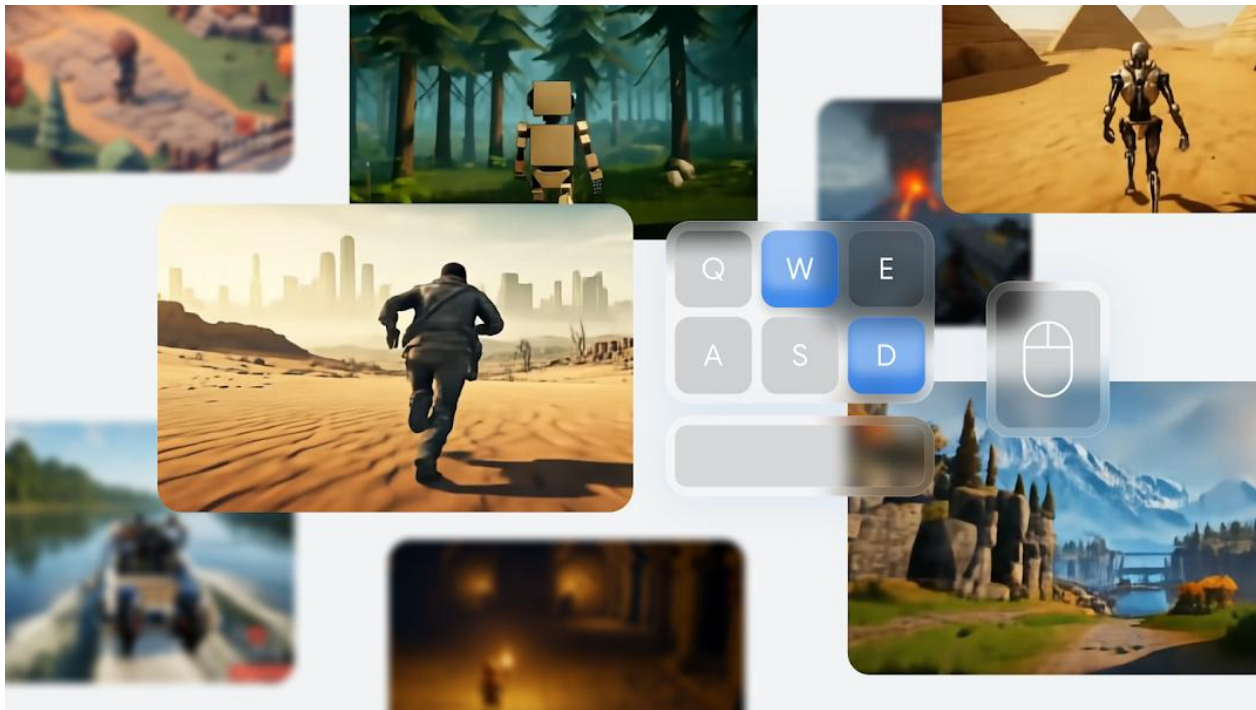
Large Video Models (LVM)



# Large Language Models

Large Video Models (LVM)

Google Genie 2



Source: <https://deepmind.google/discover/blog/genie-2-a-large-scale-foundation-world-model/>

# Tokenization

## Introduction

- process of breaking down a sequence of text into individual units
- typical units: words, subwords
- units called tokens
- different approaches
  - word tokenization
  - sentence tokenization
  - subword tokenization

# Tokenization

## Word Tokenization

### Sample Text

The quick brown fox jumps over the lazy dog.

### Tokens

The quick brown fox jumps over the lazy dog.

# Tokenization

## Sentence Tokenization

### Sample Text

The quick brown fox jumps over the lazy dog.

### Tokens

The quick brown fox jumps over the lazy dog.



# Tokenization

## Word Tokenization and Embedding

- fundamental step in NLP (Natural Language Processing)
- first step of all NLP tasks

Text

The quick brown fox jumps over the lazy dog.

Tokens

The quick brown fox jumps over the lazy dog.

Embeddings

[0.2, ...]

...

# Tokenization

## Sentence Tokenization and Embedding

- fundamental step in NLP
- first step of all NLP tasks

Text

The quick brown fox jumps over the lazy dog.

Tokens

The quick brown fox jumps over the lazy dog.

Embeddings

[0.2, ...]

...

# Tokenization

## Sub-word Tokenization

Text

It is raining.



Tokens

It

is

rain

ing

.